

Module 32: CNNs in Pytorch

```
import torch  
import matplotlib.pyplot as plt  
import numpy as np.
```

① Data loading → Loading the CIFAR10 dataset using torchvision

```
import torchvision  
import torchvision.transforms as transforms → store address  
trainset = torchvision.datasets.CIFAR10 (root = '/data',  
                                         train = True,  
                                         download = True,  
                                         transform = transforms.ToTensor())
```

transforms : applying fns at the time of loading itself

↓

Here, we are converting them to tensor type

```
classes = ('plane', 'car', 'bird', 'cat', 'deer', 'dog',  
           'frog', 'horse', 'ship', 'truck')
```

```
trainloader = torch.utils.data.DataLoader (trainset,  
                                         batch_size = 4,  
                                         shuffle = True)
```

dataiter = iter (trainloader) → way of working with
trainloader is to make an iterator on it

```
images, labels = dataiter.next()
```

```
print(images.shape) # [4, 3, 32, 32]
```

② visualize data : To plot image, we convert them into numpy.

```
img = images[0]
```

this is the order required to plot the image.

```
npimg = img.numpy()
```



```
npimg = npimg.transpose(npimg, (1, 2, 0))
```



for high order matrix, transpose means any permutation of the different axes

```
plt.figure(figsize=(1,1))
```

```
plt.imshow(npimg)
```

```
plt.show()
```



move this to a function



```
def imshow(img):
```

```
    npimg = img.numpy()
```

```
    plt.imshow(npimg.transpose(npimg, (1, 2, 0)))
```

```
    plt.show()
```

converts into a repr, where the images
get a grid repr.

```
imshow(torchvision.utils.make_grid(images))
```

```
print(' '.join(classes[labels[j]] for j in range(4)))
```

③ CNNs in Pytorch

1. single convolutional layer

class firstCNN (nn.Module) :

def __init__(self):

super(firstCNN, self).__init__()

self.conv1 = nn.Conv2d(3, 16, 3)

def forward(self, x):

return self.conv1(x)

net = firstCNN()

out = net(images)

out.shape → [4, 16, 30, 30])

i/p channels → tuple for rect.

kernel size

o/p channels .

to include padding, stride

padding = (1, 1)

stride = (2, 2) ↗ Y

X

* accessing parameters of model.

```
for param in net.parameters():
    print(param.shape)
```

torch.size([16, 3, 3, 3])

The parameters associated
with one kernel along i/p
depth

The no. of output kernels .

torch.size([16])

→ The bias

* plotting representation corresponding to first kernel
in the convolution.

```
out1 = out[0, 0, :, :].detach().numpy()  
print(out1.shape)
```

2. Deep Convolutional Network

```
class firstCNN_v2(nn.Module):  
    def __init__(self):  
        super(firstCNN_v2, self).__init__()  
        self.model = nn.Sequential(  
            nn.Conv2D(3, 8, 3),  
            nn.Conv2D(8, 16, 3)  
        )
```

$(N, 3, 32, 32) \rightarrow (N, 8, 30, 30)$

$(N, 8, 30, 30) \rightarrow (N, 16, 28, 28)$

```
def forward(self, x):  
    return self.model(x)
```

```
net = firstCNN_v2()
```

```
out = net(imager)
```

```
out.shape
```

```
plt.imshow(out[0, 0, :, :].detach().numpy())
```

class firstCNN_v3(nn.Module): 32 - 5 + 1

```
def __init__(self):  
    super(firstCNN_v3, self).__init__()  
    self.model = nn.Sequential(  
        nn.Conv2d(3, 6, 5),  
        nn.AvgPool2d(2, stride=2),  
        nn.Conv2d(6, 16, 5),  
        nn.AvgPool2d(2, stride=2)  
    )
```

mask of 2x2
→ avoids any overlap.

(N, 3, 32, 32) → (N, 6, 28, 28)

(N, 6, 28, 28) → (N, 6, 14, 14)

(N, 6, 14, 14) → (N, 16, 10, 10)

(N, 6, 10, 10) → (N, 6, 5, 5)

④ LeNet

```
class LeNet(nn.Module):  
    def __init__(self):  
        super(LeNet, self).__init__()  
        self.cnn_model = nn.Sequential(  
            nn.Conv2d(3, 6, 5)  
            nn.Tanh()  
            nn.AvgPool2d(2, stride=2)  
            nn.Conv2d(6, 16, 5)  
            nn.Tanh()  
            nn.AvgPool2d(2, stride=2)  
        )
```

self.fc_model = nn.Sequential(

nn.Linear(400, 120),

nn.Tanh()

good practice to
mention shape
modifications.

nn.Linear(120, 84),

nn.Tanh()

nn.Linear(84, 10)

)

def forward(self, x):

print(x.shape)

x = self.cnn_model(x)

print(x.shape)

batch dim kept same

x = x.view(x.size(0), -1) → flattening operation

print(x.shape)

x = self.fc_model(x)

print(x.shape)

return x

→ softmax not used, already
taken into consideration by
cross entropy loss for training

max_values, pred_class = torch.max(out.data, 1)

print(pred_class)

⑤ Training Lenet

batch_size = 128

trainset = torchvision.datasets.CIFAR10(root='./data',
train=True,
download=True,

```
transform = transforms.ToTensor()
```

```
trainloader = torch.utils.data.DataLoader(trainset,  
                                         batch_size=batch_size,  
                                         shuffle=True)
```

```
testset = torchvision.datasets.CIFAR10(root='./data',  
                                       train=False,  
                                       download=True,  
                                       transform=transforms.ToTensor())
```

```
testloader = torch.utils.data.DataLoader(testset,  
                                         batch_size=batch_size,  
                                         shuffle=False)
```

```
def evaluation(dataloader):  
    total, correct = 0, 0  
    for data in dataloader:  
        inputs, labels = data  
        outputs = net(inputs)  
        pred = torch.max(outputs.data, 1)  
        total += labels.size(0)  
        correct += (pred == labels).sum().item()  
  
    return correct / total * 100
```

```
net = LeNet()
```

```
import torch.optim as optim  
loss_fn = nn.CrossEntropyLoss()  
opt = optim.Adam(net.parameters())  
  
*) Learning loop:  
    /.\ time  
    loss_arr = []  
    loss_epoch_arr = []  
    max_epochs = 16  
  
    for epoch in range(max_epochs):  
        for i, data in enumerate(trainloader, 0):  
            inputs, labels = data  
            opt.zero_grad()  
            outputs = net(inputs)  
            loss = loss_fn(outputs, labels)  
            loss.backward()  
            opt.step()  
  
            loss_arr.append(loss.item())  
  
        loss_epoch_arr.append(loss.item())  
  
        print(f'Epoch {epoch}/{max_epochs}, evaluation({testloader}), evaluation({trainloader})')
```

```
plt.plot (loss_epoch_an)
```

```
plt.show()
```

⑥ Move to GPU → Cnet, i/p, o/p on GPU

```
device = torch.device ("cuda:0" if torch.cuda.is_available ()  
else "cpu")
```

```
def evaluation (dataloader):
```

```
    total, correct = 0, 0
```

```
    for data in dataloader :
```

```
        inputs, labels = data
```

```
        inputs, labels = inputs.to (device), labels.to (device)
```

```
        outputs = net (inputs)
```

```
- , pred = torch.max (outputs.data, 1)
```

```
        total += labels.size (0)
```

```
        correct += (pred == labels).sum ().item ()
```

```
    return correct / total * 100
```

```
net = LeNet (), to (device)
```

```
loss_fn = nn. CrossEntropyLoss ()
```

```
opt = optim. Adam (net. parameters ())
```

%.%. time

max_epochs = 16

```
for epoch in range (max_epochs) :
```

```
    for i, data in enumerate (trainloader, 0) :
```

```
        inputs, labels = data
```

```
inputs, labels = inputs.to(device), labels.to(device)
opt.zero_grad()
outputs = net(inputs)
loss = loss_fn(outputs, labels)
loss.backward()
opt.step()

loss_avg.append(loss.item())
loss_epoch_avg.append(loss.item())

print(f'Epoch {epoch}/{max_epochs}, evaluation ({testloader}), evaluation ({trainloader})
```

⑦ Basic visualisation

```
imshow(torchvision.utils.make_grid(images))
```

net = net.to("cpu")

out = net(images) → gives the final output

out = net.cnn_model[0](images) → sending images through
out.shape sequence of layers that layer
getting the first layer

image_id = 3

```
plt.figure(figsize=(2, 2))
```

```
imshow(images[image_id])
```

```
plt.figure(figsize = (6,6))
plt.subplot(321)
for i in range(6):
    ax1 = plt.subplot(3,2, i+1)
    plt.imshow(out[image_id, i, :, :].detach().numpy(), cmap = "binary")
```

```
plt.show()
```

⑧ Exercises

- i) Tank to ReLU in LeNet
- ii) Overfitting observed → use weight-decay in torch.optim to add L2 regularisation.
- iii) move LeNet to MNIST from CIFAR10.